

# Introduction to Reinforcement Learning with Function Approximation

Norio Kosaka

January 26, 2019

## Contents

<b>1</b>	<b>Foreword</b>	<b>2</b>
<b>2</b>	<b>Introduction to RL: Successes and Challenged</b>	<b>2</b>
2.1	What is RL? . . . . .	2
2.2	Some RL Successes . . . . .	2
2.2.1	DQN . . . . .	3
<b>3</b>	<b>The Formal Problem: Finite Markov Decision Processes</b>	<b>3</b>
3.1	The Environment: A Finite Markov Decision Process(MDP) . . .	3
<b>4</b>	<b>Exact Solution Methods (tabular methods)</b>	<b>4</b>
4.1	Q-Learning, the simplest model-free RL algorithm . . . . .	4
4.2	Policy Improvement Theorem . . . . .	4
4.3	Policy Iteration . . . . .	4
4.4	Exploration/Exploitation Dilemma . . . . .	5
4.5	Bootstrapping . . . . .	5
4.6	Summary . . . . .	5
<b>5</b>	<b>Approximate Solution Methods (function approximation)</b>	<b>5</b>
5.1	Function Approximations . . . . .	6
5.2	Problem of instability with semi-gradient Q-learning . . . . .	6
5.3	Workaround for Deadly Triad . . . . .	7
<b>6</b>	<b>Miscellany and closing remarks</b>	<b>7</b>
6.1	Many dimensions of RL . . . . .	7
6.1.1	Problems . . . . .	7
6.1.2	Methods . . . . .	7
6.2	Policy-gradient actor-critic methods . . . . .	8

# 1 Foreword

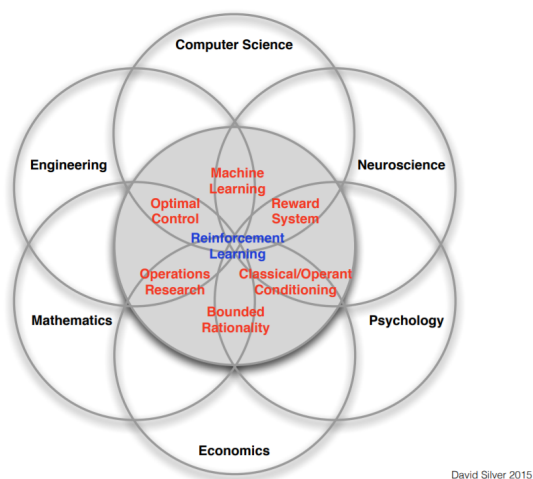
This is the summary of the great presentation by Richard Sutton's presentation at NIPS 2015. And the original PPT is here:

<http://media.nips.cc/Conferences/2015/tutorialslides/SuttonIntroRL-nips-2015-tutorial.pdf>

## 2 Introduction to RL: Successes and Challenged

### 2.1 What is RL?

- Agent-oriented learning - learning by interacting with an environment achieve a goal
- Learning by trial and error, with only delayed evaluative feedback(reward)
- The beginnings of a science of mind that is neither natural science nor applications technology



### 2.2 Some RL Successes

- Learned the world's best player of Backgammon(Tesauro 1995 [14])
- Learned acrobatic helicopter autopilots(Ng, Abbeel, Coates et al., 2006+ [1])
- Widely used in the placement and selection of advertisements and pages on the web (e.g., A-B tests)
- Used to make strategic decisions in *Jeopardy!*(IBM's Watson 2011)

- Achieved human-level performance on Atari games from pixel-level visual input, in conjunction with deep learning (Google Deepmind 2015 [10], [3])

### 2.2.1 DQN

- Learned to play 49 games for the Atari 2600 game console, without labels or human input, from self-play and the score alone
- Learned to play better than all previous algorithms and at human level for more than half the games
- mapping raw screen pixels to predictions of final score for each of 18 joystick actions: Same learning algorithm applied to all 49 game without human tuning!!

## 3 The Formal Problem: Finite Markov Decision Processes

### 3.1 The Environment: A Finite Markov Decision Process(MDP)

- Proposed by Howard, 1964 [6]
- Discrete time:  $t = 1, \dots, T$
- A finite set of **states**
- A finite set of **actions**
- A finite set of **rewards**
- Life is a trajectory:  $S_0, A_0, S_1, A_1, \dots, S_t, A_t$
- with arbitrary Markov(stochastic, state-dependent) dynamics:

$$p(r, s' | s, a) = P[R_{t+1} = r, S_{t+1} = s' | S_t = s, A_t = a]$$

- **Policies:** deterministic policy:  $a = \pi(s)$
- Agent learns to find the best policy which maximise the return
- **Action-value functions:** it says how good it is to be in a state, take an action, and thereafter follow a policy

$$q(s, a) = E[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots | S_t = s, A_t = a, A_{t+1:\infty} \sim \pi]$$

- **Optimal Policies( $\pi^*$ ):** maximise the action-value function:

$$q_{\pi^*}(s, a) = \max_{\pi} q_{\pi}(s, a) = q_*(s, a)$$

$$\pi^*(s) = \arg \max_a q_*(s, a)$$

## 4 Exact Solution Methods (tabular methods)

### 4.1 Q-Learning, the simplest model-free RL algorithm

It is introduced by Watkins&Dayan 1992 [16]

- Initialise an array  $Q(s, a)$  arbitrarily
- Choose actions in any way, perhaps based on  $Q$ , such that all actions are taken in all states
- On each time step, change one element of the array

$$\Delta Q(s, a) = \alpha \left( R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t) \right)$$

- If desired, reduce the step-size parameter  $\alpha$  over time

### 4.2 Policy Improvement Theorem

Given the value function for any policy  $\pi$ :  $q_\pi(s, a) \forall s, a$ , it can always be greedified to obtain a better policy


$$\pi'(s) = \arg \max_a q_\pi(s, a)$$

where better means

$$q_{\pi'}(s, a) \geq q_\pi(s, a) \forall s, a$$

### 4.3 Policy Iteration

- Policy Evaluation
- Policy Improvement



## 4.4 Exploration/Exploitation Dilemma

- You can't do the action that you think is best all the time
- You can't explore all the time
- Thus you must both explore and exploit, but neither to excess. What is the right balance?

## 4.5 Bootstrapping

Bootstrapping is the key idea underlying both **dynamic programming (DP)** and all **temporal-difference (TD)** learning. It updates an estimate from an estimate, it is like a guess from a guess. The concept itself is derived from the **Bellman Expectation Equation**:

$$\begin{aligned} q_{\pi}(s, a) &= \mathbb{E} \left[ R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots | S_t = s, A_t = a, A_{t+1:\infty} \sim \pi \right] \\ &= \mathbb{E} \left[ R_{t+1} + \gamma q_{\pi}(S_{t+1}, A_{t+1} | S_t = s, A_t = a, A_{t+1} \sim \pi) \right] \end{aligned}$$

Or the Bellman optimality equation:

$$q_*(s, a) = \mathbb{E} \left[ R_{t+1} + \gamma \max_{a'} q_*(S_{t+1}, a') | S_t = s, A_t = a \right]$$

## 4.6 Summary

**Off-policy** learning is learning about the value of a policy other than the policy being used to generate the trajectory. Q-learning learns about the value of its deterministic greedy policy—which gradually become optimal—from data while behaving in a more exploratory manner. Thus Q-learning is off-policy and this is essential to its strategy for escaping the **explore/exploit** dilemma

- the target policy is the policy being learned about
- the behavior policy is the policy generating the trajectory data
- on-policy learning is when the two policies are the same

## 5 Approximate Solution Methods (function approximation)

So far, we have confirmed that RL finds optimal policies for finite state-space environment, if the value functions and policies can be exactly represented in tables. But the real world is too large and complex for tables; continuous state-space. So, the question we will look at in this section is;

- Will RL work with approximations?
- Will RL work with function approximators?

## 5.1 Function Approximations

- Represent the action-value function by a **parameterised function approximator** with parameter  $\theta$

$$q(s, a, \theta) \approx q_*(s, a)$$

- The approximator could be a **deep neural network**, with the weights being the parameter

Q-learning works well with function approximation(Watkins 1989 [17])

- Semi-gradient Q-learning(Watkins 1989 [17])

$$L(\theta) = \mathbb{E} \left[ \left( R_{t+1} + \gamma \max_a \hat{q}(S_{t+1}, a, \theta) - \hat{q}(S_t, A_t, \theta) \right)^2 \right] \quad (1)$$

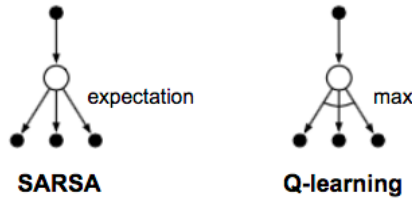
$$\Delta \theta_t = \alpha \left( R_{t+1} + \gamma \max_a \hat{q}(S_{t+1}, a, \theta) - \hat{q}(S_t, A_t, \theta) \right) \frac{\partial \hat{q}(S_t, A_t, \theta)}{\partial \theta_t} \quad (2)$$

- Semi-gradient Sarsa(Rummery 1994 [11], Sutton 1988 [12])

$$L(\theta) = \mathbb{E} \left[ \left( R_{t+1} + \gamma \hat{q}(S_{t+1}, a, \theta) - \hat{q}(S_t, A_t, \theta) \right)^2 \right] \quad (3)$$

$$\Delta \theta_t = \alpha \left( R_{t+1} + \gamma \hat{q}(S_{t+1}, a, \theta) - \hat{q}(S_t, A_t, \theta) \right) \frac{\partial \hat{q}(S_t, A_t, \theta)}{\partial \theta_t} \quad (4)$$

Figure 1: Backup: SARSA VS Q-LEARNING



## 5.2 Problem of instability with semi-gradient Q-learning

As an on-policy method, Semi-gradient Sarsa has good convergence properties(Sutton 1988 [12], Dayan 1992 [4], Tsitsiklis&Van Roy 1997 [15]) But when it comes to Q-learning, the risk of divergence arises whenever we combine three things below;

- Function approximation: significantly generalizing from large numbers of examples

- Bootstrapping: learning value estimates from other value estimates, as in dynamic programming and temporal-difference learning. but it introduces **biases** in learning
- Off-policy learning: learning about a policy from data not due to that policy, as in Q-learning, where we learn about the greedy policy from data with a necessarily more exploratory policy

So, this problem is called **The deadly triad**.

### 5.3 Workaround for Deadly Triad

- Tsitsiklis&Van Roy 1997 [15] proposed: the degree of bootstrapping:  $\lambda$ , from  $\lambda = 0$ (full bootstrapping) to  $\lambda = 1$ (no bootstrapping)
- Double Q-learning, van Hasselt 2010 [5], but it is too soon to be sure
- least-squares methods like off-policy LSTD( $\lambda$ ), (Yu 2010 [18], Mahmood et al. 2015 [9]), but their computational costs scale with the square of the number of parameters
- Gradient-TD (Maei, 2011 [8]) and proximal-gradientTD (Mahadevan et al., 2015 [7]), These seem to me(R.Sutton) to be the best attempts to make TD methods with the robust convergence properties of stochastic gradient descent.
- Residual gradient methods (Baird 1999 [2])
- Emphatic-TD methods (Sutton, White&Mahmood 2015 [13], Yu 2015 [19])

## 6 Miscellany and closing remarks

### 6.1 Many dimensions of RL

#### 6.1.1 Problems

- prediction vs control
- MDPs vs Bandits

#### 6.1.2 Methods

- Tabular vs function approximation
- On-policy vs off-policy
- Bootstrapping vs Monte Carlo
- Model-based vs Model-free
- value-based vs policy-based

## 6.2 Policy-gradient actor-critic methods

- Policy is explicitly represented with its own parameters independent of any value function, so it is easy and simpler than dealing with the value function
- Policy parameters are updated by stochastic gradient ascent in a performance measure such as average reward per step
- A state-value function (critic) is optional but can significantly reduce variance
- Good convergence properties (on-policy)

## References

- [1] Pieter Abbeel, Adam Coates, and Andrew Y Ng. “Autonomous helicopter aerobatics through apprenticeship learning”. In: *The International Journal of Robotics Research* 29.13 (2010), pp. 1608–1639.
- [2] Leemon Baird. “Residual algorithms: Reinforcement learning with function approximation”. In: *Machine Learning Proceedings 1995*. Elsevier, 1995, pp. 30–37.
- [3] Marc G Bellemare et al. “The arcade learning environment: An evaluation platform for general agents”. In: *Journal of Artificial Intelligence Research* 47 (2013), pp. 253–279.
- [4] Peter Dayan. “The convergence of TD ( $\lambda$ ) for general  $\lambda$ ”. In: *Machine learning* 8.3-4 (1992), pp. 341–362.
- [5] Hado V Hasselt. “Double Q-learning”. In: *Advances in Neural Information Processing Systems*. 2010, pp. 2613–2621.
- [6] Ronald A Howard. “Dynamic programming and Markov processes”. In: (1964).
- [7] Bo Liu et al. “Finite-Sample Analysis of Proximal Gradient TD Algorithms.” In: *UAI*. Citeseer. 2015, pp. 504–513.
- [8] Hamid Reza Maei and Richard S Sutton. *Gradient temporal-difference learning algorithms*. University of Alberta Edmonton, Alberta, 2011.
- [9] Ashique Rupam Mahmood and Richard S Sutton. “Off-policy learning based on weighted importance sampling with linear computational complexity.” In: *UAI*. Citeseer. 2015, pp. 552–561.
- [10] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (2015), p. 529.
- [11] Gavin A Rummery and Mahesan Niranjana. *On-line Q-learning using connectionist systems*. Vol. 37. University of Cambridge, Department of Engineering Cambridge, England, 1994.



- [12] Richard S Sutton. “Learning to predict by the methods of temporal differences”. In: *Machine learning* 3.1 (1988), pp. 9–44.
- [13] Richard S Sutton, A Rupam Mahmood, and Martha White. “An emphatic approach to the problem of off-policy temporal-difference learning”. In: *The Journal of Machine Learning Research* 17.1 (2016), pp. 2603–2631.
- [14] Gerald Tesauro. “Temporal difference learning and TD-Gammon”. In: *Communications of the ACM* 38.3 (1995), pp. 58–68.
- [15] JN Tsitsiklis and B Van Roy. *An analysis of temporal-difference learning with function approximation* Technical. Tech. rep. Report LIDS-P-2322). Laboratory for Information and Decision Systems . . . , 1996.
- [16] Christopher JCH Watkins and Peter Dayan. “Q-learning”. In: *Machine learning* 8.3-4 (1992), pp. 279–292.
- [17] Christopher John Cornish Hellaby Watkins. “Learning from delayed rewards”. PhD thesis. King’s College, Cambridge, 1989.
- [18] Huizhen Yu. “Convergence of least squares temporal difference methods under general conditions”. In: *Proceedings of the 27th International Conference on Machine Learning (ICML-10)*. Citeseer. 2010, pp. 1207–1214.
- [19] Huizhen Yu. “On convergence of emphatic temporal-difference learning”. In: *Conference on Learning Theory*. 2015, pp. 1724–1751.